

MLNX-OS Up/Downgrade via Ansible v1.0

HyungKwang Choi (2021.9)

Table of Contents

1.	Test environment.....	3
1.1	Topology.....	3
1.2	Device info.....	3
1.2.1	Server	3
1.2.1	SB7700.....	3
1.3	Scripts.....	5
1.4	License.....	5
1.4.1	if your MLNX-OS version is above v3.7, you do not need a license.	5
1.4.2	if your MLNX-OS version is lower than v3.7, you need a license.	5
2.	Running.....	7
2.1	Introduction on my Ansible script.....	7
2.2.1	The Purpose.....	7
2.2.2	What it does	7
2.2.3	Exact steps.....	8
2.2.4	How to run continuously.....	9
2.2	How to run	9
2.3	Running result : “Output.text”	9
3.	Possible script failures	11
3.1	Ansible & Python version mis-match.....	11
3.2	License invalid.....	11
4.	Reference	12
4.1	Ansible module for Mellanox	12
4.2	From NVIDIA Doc, how to	12

1. Test environment

1.1 Topology.

CentOS Server (ansible server)-----remote-----standalone SB7700

1.2 Device info.

1.2.1 Server

- CentOS 8.0
- Ansible 2.9.24
- Python 3.6.8
- Please place images, which you are running OS Up/Downgrade.

```
[root@NVIDIA ~]# cat /etc/redhat-release
CentOS Linux release 8.0.1905 (Core)

[root@NVIDIA ~]# ansible --version
ansible 2.9.24
  config file = /etc/ansible/ansible.cfg
  configured module search path = ['/root/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']
  ansible python module location = /usr/lib/python3.6/site-packages/ansible
  executable location = /usr/bin/ansible
  python version = 3.6.8 (default, May 21 2019, 23:51:36) [GCC 8.2.1 20180905 (Red Hat 8.2.1-3)]

[root@NVIDIA ~]# ls -al | grep image
-rw-r--r-- 1 root root 442345149 Aug 24 02:21 image-X86_64-3.6.2102.img
-rw-r--r-- 1 root root 591062588 Aug 24 02:22 image-X86_64-3.6.6162.img
-rw-r--r-- 1 root root 599453042 Aug 23 16:29 image-X86_64-3.6.8008.img
-rw-r--r-- 1 root root 599644444 Aug 23 14:24 image-X86_64-3.6.8010.img
```

1.2.1 SB7700

```
NVIDIA [standalone: master] (config) # show version
Product name:  MLNX-OS
Product release:  3.6.8010
Build ID:      #1-dev
Build date:    2018-08-20 18:04:19
Target arch:   x86_64
```

Target hw: x86_64
Built by: jenkins@XXXXXXX
Version summary: X86_64 3.6.8010 2018-08-20 18:04:19 x86_64

NVIDIA [standalone: master] (config) # [show inventory](#)

Module	Part Number	Serial Number	Asic Rev.	HW Rev.
CHASSIS	MSB7700-ES2F	MTXXXXXXX	N/A	A7
MGMT	MSB7700-ES2F	MTXXXXXXX	0	A7
FAN1	MTEF-FANF-A	MTXXXXXXX	N/A	A3
FAN2	MTEF-FANF-A	MTXXXXXXX	N/A	A3
FAN3	MTEF-FANF-A	MTXXXXXXX	N/A	A3
FAN4	MTEF-FANF-A	MTXXXXXXX	N/A	A3
PS1	MTEF-PSF-AC-A	MTXXXXXXX	N/A	A2
PS2	MTEF-PSF-AC-A	MTXXXXXXX	N/A	A2

NVIDIA [standalone: master] (config) # [show images](#)

Installed images:

Partition 1:

version: X86_64 3.6.6162 2018-08-03 10:36:36 x86_64

Partition 2:

version: X86_64 3.6.8010 2018-08-20 18:04:19 x86_64

Last boot partition: 2

Next boot partition: 2

Images available to be installed:

1:

Image : image-X86_64-3.6.8010.img

Version: X86_64 3.6.8010 2018-08-20 18:04:19 x86_64

Serve image files via HTTP/HTTPS: no

No image install currently in progress.

Boot manager password is set.

Image signing : trusted signature always required

Admin require signed images: yes

Settings for next boot only:

Fallback reboot on configuration failure: yes (default)

1.3 Scripts

1. Ansible “OS_Up_downgrade.yml”
2. Inventory_hostname file

```
[root@NVIDIA ~]# cat /etc/ansible/hosts
[ONYX]
NVIDIA.test.labs ◀==== modify properly.

[ONYX:vars]
ansible_network_os=onyx
ansible_become=yes
ansible_become_method=enable
ansible_ssh_user=admin
ansible_ssh_pass=admin
```

1.4 License

- License is not mandatory.
- If you want to check the SSD status or “Erase Count”, you need a license if version is lower than 3.7

1.4.1 if your MLNX-OS version is above v3.7, you do not need a license.
Because the command “#(config)fae show smart” is applicable as of v3.7.

1.4.2 if your MLNX-OS version is lower than v3.7, you need a license.
Because “#(config)fae show smart” does not work below v3.7
Then you need a license to check SSD status.

1). How to get license. You need mgmt mac-address and contact CSI lab admin

```
NVIDIA [standalone: master] (config) # show interfaces mgmt0 ◀=====
Interface mgmt0 status:
  Comment:
  Admin up:      yes
  Link up:       yes
  DHCP running:  yes
  IP address:    XXXX.XXXX.XXXX.XXXX
  Netmask:      255.255.252.0
  IPv6 enabled:  yes
  Autoconf enabled: no
  Autoconf route: yes
  Autoconf privacy: no
  DHCPv6 running: yes (but no valid lease)
```

IPv6 addresses: 1
IPv6 address: XXXX::XXXX:XXXX:XXXX:XXXX/64
Speed: 1000Mb/s (auto)
Duplex: full (auto)
Interface type: ethernet
Interface source: bridge
MTU: 1500
HW address: XX:XX:XX:XX:XX:XX ←

2). How to see SSD status or “Erase Count”

```
(config) # license install <license key>
(config) # show licenses
(config) # exec iSmart -d /dev/sda
```

NVIDIA [standalone: master] # show licenses

NVIDIA [standalone: master] # exec iSmart -d /dev/sda

```
*****
* Innodisk iSMART V3.9.41                               2018/05/25 *
*****
```

Model Name: InnoDisk Corp. - mSATA 3ME
FW Version: S141119O
Serial Number: XXXXXXXXXXXXXXXXXXXXX
Health: 0.00
Capacity: 14.914146 GB
P/E Cycle: 3000
Lifespan : 0 (Years : 0 Months : 0 Days : 0)
Write Protect: Disable
InnoRobust: Disable

ID	SMART Attributes	Value	Raw Value
[09]	Power On Hours	[53455]	[0900000000CFD00000000000]
[0C]	Power Cycle Count	[104]	[0C0000000068000000000000]
[AA]	Total Bad Block Count	[13]	[AA0300646400000D00000000]
[AD]	Erase Count Max.	[5861]	[AD12006464DE15E516000000]
[AD]	Erase Count Avg.	[5598]	[AD12006464DE15E516000000]
[C2]	Temperature	[0]	[000000000000000000000000]
[EB]	Later Bad Block	[0]	[EB0200640000000000000000]
[EB]	Read Block	[0]	[EB0200640000000000000000]
[EB]	Write Block	[0]	[EB0200640000000000000000]
[EB]	Erase Block	[0]	[EB0200640000000000000000]
[EC]	Unstable Power Count	[0]	[EC0200646400000000000000]

```
(config) # license delete 1
```

2. Running

2.1 Introduction on my Ansible script

2.2.1 The Purpose

- It does MLNX-OS upgrade and downgrade continuously for replication test purpose.

While I am handling cases, sometimes, there are some customer complaints about “during OS/FW upgrade or downgrade, SSD gets stuck or errored block, Erase count increases”.

To verify that we TAC engineers sets up a lab. So to run it automatically via Ansible, we can save time and focus on the SSD data collected during OS upgrade/downgrade.

2.2.2 What it does

1). Doing OS upgrade -> downgrade -> upgrade.....

3.6.2102 -> 3.6.6162 -> 3.6.8010 -> 3.6.6162 -> 3.6.2102.....

2). During that, it writes current version, collects SSD status and write to a output.txt file

```
# print/write the current version
- name: print_current_version
  debug:
    msg: "Current version {{ current_version }}"
- name: Write Version
  local_action: shell echo "Current version {{ current_version }}" >> /root/output.txt

# Add license
- name: Add license
  onyx_config:
    lines:
      - license install XXXX_XXXXXX_XXXXXXX_XXXXXXX_XXXXXX_XXXXXXX
  register: add_license

# Get the SSD status
- name: _exec iSmart -d /dev/sda
  onyx_command:
    commands: _exec iSmart -d /dev/sda
  register: ssd_data
```

```
# Write the SSD status to output.txt file
- name: SSD status
  local_action: shell echo "{{ ssd_data.stdout }}" >> /root/output.txt
```

2.2.3 Exact steps

```
# Runing main task
- name: onyx_upgrade
~
~
tasks: ←=== main task

# get current switch Onyx version with command "show version concise" and save it to a variable named curr_version
- name: get_current_switch_version

# save the onyx version as a fact named current_version
- name: set_fact_current_version

# print/write the current version
- name: print_current_version

# Add license
- name: Add license

# Get the SSD status
- name: _exec iSmart -d /dev/sda

# Write the SSD status to output.txt file
- name: SSD status

# run the delete images command
- name: delete images

# check if the img file exists in location
- name: check_source_file_existence

# Scp file copy
- name: copy_source_file_to_device

# Image install
- name: image install

# # image boot next
- name: image boot next

# # save config
- name: save config

# run the reload command
- name: reload
```



```
# use wait_for to wait for port 22 to be available on the switch after reboot
- name: wait for switch to reboot
```

2.2.4 How to run continuously.

Above “Exact steps” is 1 cycle. So it repeated only changing version info. So if you want to run many cycle, just copy & paste whole steps as many as you want to run.

2.2 How to run

```
[root@NVIDIA ~]# ansible-playbook OS_Up_downgrade.yml
```

➤ For debugging

```
[root@NVIDIA ~]# ansible-playbook OS_Up_downgrade.yml -vvvv
```

2.3 Running result : “Output.text”

```
[root@NVIDIA ~]# ansible-playbook OS_Up_downgrade.yml
```

```
[root@NVIDIA ~]# cat output.txt
```

Current version 3.6.2102

* Innodisk iSMART V3.9.41 2018/05/25 *

Model Name: InnoDisk Corp. - mSATA 3ME

FW Version: S141119O

Serial Number: XXXXXXXXXXXXXXXXXXXX

Health: 0.00

Capacity: 14.914146 GB

P/E Cycle: 3000

Lifespan : 0 (Years : 0 Months : 0 Days : 0)

Write Protect: Disable

InnoRobust: Disable

ID	SMART Attributes	Value	Raw Value
[09]	Power On Hours	[53455]	[0900000000CFD00000000000]

[0C] Power Cycle Count	[104]	[0C0000000068000000000000]
[AA] Total Bad Block Count	[13]	[AA03006464000000D000000000]
[AD] Erase Count Max.	[5861]	[AD12006464DE15E516000000]
[AD] Erase Count Avg.	[5598]	[AD12006464DE15E516000000]
[C2] Temperature	[0]	[000000000000000000000000]
[EB] Later Bad Block	[0]	[EB0200640000000000000000]
[EB] Read Block	[0]	[EB0200640000000000000000]
[EB] Write Block	[0]	[EB0200640000000000000000]
[EB] Erase Block	[0]	[EB0200640000000000000000]
[EC] Unstable Power Count	[0]	[EC0200646400000000000000]

3. Possible script failures

3.1 Ansible & Python version mis-match.

➤ My lab environment. Please check and compare it with yours

- CentOS 8.0
- Ansible 2.9.24
- Python 3.6.8

3.2 License invalid.

➤ One of my script running is “Adding license & SSD status.”. Please modify/add license properly.

```
# Add license

- name: Add license
  onyx_config:
    lines:
      - license install XXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX
  register: add_license

# Get the SSD status

- name: _exec iSmart -d /dev/sda
  onyx_command:
    commands: _exec iSmart -d /dev/sda
  register: ssd_data
```

1). My script is written based on v3.6. So for this, license required to check SSD status.

2). If you want to skip to check SSD status, simply remove the above statement.

```
NVIDIA [standalone: master] # show licenses
License 1: XXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX
Feature:    RESTRICTED_CMDS_GEN2
Description: Access to restricted system functionality
Valid:      yes
End date:   2023/07/13 (ok)
Tied to MAC addr: XX:XX:XX:XX:XX:XX (ok)
Active:     yes ←=====
```

4. Reference

4.1 Ansible module for Mellanox

https://docs.ansible.com/ansible/latest/collections/mellanox/onyx/onyx_command_module.html

4.2 From NVIDIA Doc, how to

<https://docs.mellanox.com/display/ONYXv381208/Ansible>

<https://community.mellanox.com/s/article/Network-Automation-with-ZTP-and-Ansible-ONYX>