

MLNX-OS Up/Downgrade via Ansible v2.0

HyungKwang Choi (2021.9)

Table of Contents

1.	Test environment.....	3
1.1	Topology.....	3
1.2	Device info.....	3
1.2.1	Server	3
1.2.1	SB7700.....	3
1.3	Scripts.....	5
1.4	License.....	5
1.4.1	if your MLNX-OS version is above v3.7, you do not need a license.	5
1.4.2	if your MLNX-OS version is lower than v3.7, you need a license.	5
2.	Running.....	7
2.1	Introduction on my Ansible script.....	7
2.2.1	The Purpose.....	7
2.2.2	What it does	7
2.2.3	Exact steps.....	8
2.2.4	How to run continuously.....	9
2.2	How to run	9
2.3	Running result : “Output.text”	9
3.	Possible script failures	11
3.1	Ansible & Python version mis-match.....	11
3.2	License invalid.....	11
4.	Reference	12
4.1	Ansible module for Mellanox	12
4.2	From NVIDIA Doc, how to	12
5.	Errors	13
5.1	error messages running lower version Ansible.	13
5.2	Command timeout.....	14

1. Test environment

1.1 Topology.

CentOS Server (ansible server)-----remote-----standalone SB7700

1.2 Device info.

1.2.1 Server

- CentOS 8.0
- Ansible 2.9.24
- Python 3.6.8
- Please place images, which you are running OS Up/Downgrade.

```
[root@NVIDIA ~]# cat /etc/redhat-release
CentOS Linux release 8.0.1905 (Core)

[root@NVIDIA ~]# ansible --version
ansible 2.9.24
  config file = /etc/ansible/ansible.cfg
  configured module search path = ['/root/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']
  ansible python module location = /usr/lib/python3.6/site-packages/ansible
  executable location = /usr/bin/ansible
  python version = 3.6.8 (default, May 21 2019, 23:51:36) [GCC 8.2.1 20180905 (Red Hat 8.2.1-3)]

[root@NVIDIA ~]# ls -al | grep image
-rw-r--r-- 1 root root 442345149 Aug 24 02:21 image-X86_64-3.6.2102.img
-rw-r--r-- 1 root root 591062588 Aug 24 02:22 image-X86_64-3.6.6162.img
-rw-r--r-- 1 root root 599453042 Aug 23 16:29 image-X86_64-3.6.8008.img
-rw-r--r-- 1 root root 599644444 Aug 23 14:24 image-X86_64-3.6.8010.img
```

1.2.1 SB7700

```
NVIDIA [standalone: master] (config) # show version
Product name: MLNX-OS
```

Product release: 3.6.8010
Build ID: #1-dev
Build date: 2018-08-20 18:04:19
Target arch: x86_64
Target hw: x86_64
Built by: jenkins@XXXXXXX
Version summary: X86_64 3.6.8010 2018-08-20 18:04:19 x86_64

NVIDIA [standalone: master] (config) # [show inventory](#)

Module	Part Number	Serial Number	Asic Rev.	HW Rev.
CHASSIS	MSB7700-ES2F	MTXXXXXXX	N/A	A7
MGMT	MSB7700-ES2F	MTXXXXXXX	0	A7
FAN1	MTEF-FANF-A	MTXXXXXXX	N/A	A3
FAN2	MTEF-FANF-A	MTXXXXXXX	N/A	A3
FAN3	MTEF-FANF-A	MTXXXXXXX	N/A	A3
FAN4	MTEF-FANF-A	MTXXXXXXX	N/A	A3
PS1	MTEF-PSF-AC-A	MTXXXXXXX	N/A	A2
PS2	MTEF-PSF-AC-A	MTXXXXXXX	N/A	A2

NVIDIA [standalone: master] (config) # [show images](#)

Installed images:

Partition 1:

version: X86_64 3.6.6162 2018-08-03 10:36:36 x86_64

Partition 2:

version: X86_64 3.6.8010 2018-08-20 18:04:19 x86_64

Last boot partition: 2

Next boot partition: 2

Images available to be installed:

1:

Image : image-X86_64-3.6.8010.img

Version: X86_64 3.6.8010 2018-08-20 18:04:19 x86_64

Serve image files via HTTP/HTTPS: no

No image install currently in progress.

Boot manager password is set.

Image signing : trusted signature always required

Admin require signed images: yes

Settings for next boot only:

Fallback reboot on configuration failure: yes (default)

1.3 Scripts

1. Ansible “OS_Up_downgrade.yml”
2. Inventory_hostname file

```
[root@NVIDIA ~]# cat /etc/ansible/hosts
[ONYX]
NVIDIA.test.labs ⬅==== modify properly.

[ONYX:vars]
ansible_network_os=onyx
ansible_become=yes
ansible_become_method=enable
ansible_ssh_user=admin
ansible_ssh_pass=admin
```

1.4 License

- License is not mandatory.
- If you want to check the SSD status or “Erase Count”, you need a license if version is lower than 3.7

1.4.1 if your MLNX-OS version is above v3.7, you do not need a license.
Because the command “#(config)fae show smart” is applicable as of v3.7.

1.4.2 if your MLNX-OS version is lower than v3.7, you need a license.
Because “#(config)fae show smart” does not work below v3.7
Then you need a license to check SSD status.

1). How to get license. You need mgmt mac-address and contact CSI lab admin

```
NVIDIA [standalone: master] (config) # show interfaces mgmt0 ⬅=====
Interface mgmt0 status:
  Comment:
  Admin up:      yes
  Link up:       yes
  DHCP running:  yes
  IP address:    XXXX.XXXX.XXXX.XXXX
  Netmask:       255.255.252.0
  IPv6 enabled:  yes
  Autoconf enabled: no
```

```
Autoconf route:  yes
Autoconf privacy: no
DHCPv6 running:  yes (but no valid lease)
IPv6 addresses:  1
IPv6 address:    XXXX::XXXX:XXXX:XXXX:XXXX/64
Speed:           1000Mb/s (auto)
Duplex:           full (auto)
Interface type:  ethernet
Interface source: bridge
MTU:              1500
HW address:      XX:XX:XX:XX:XX:XX ←
```

2). How to see SSD status or “Erase Count”

```
(config) # license install <license key>
(config) # show licenses
(config) # exec iSmart -d /dev/sda
```

```
NVIDIA [standalone: master] # show licenses
```

```
NVIDIA [standalone: master] # exec iSmart -d /dev/sda
```

```
*****
* Innodisk iSMART V3.9.41                               2018/05/25 *
*****
```

```
Model Name: InnoDisk Corp. - mSATA 3ME
FW Version: S141119O
Serial Number: XXXXXXXXXXXXXXXXXXXX
Health: 0.00
Capacity: 14.914146 GB
P/E Cycle: 3000
Lifespan : 0 (Years : 0 Months : 0 Days : 0)
Write Protect: Disable
InnoRobust: Disable
```

ID	SMART Attributes	Value	Raw Value
[09]	Power On Hours	[53455]	[0900000000CFD00000000000]
[0C]	Power Cycle Count	[104]	[0C0000000068000000000000]
[AA]	Total Bad Block Count	[13]	[AA0300646400000D00000000]
[AD]	Erase Count Max.	[5861]	[AD12006464DE15E516000000]
[AD]	Erase Count Avg.	[5598]	[AD12006464DE15E516000000]
[C2]	Temperature	[0]	[000000000000000000000000]
[EB]	Later Bad Block	[0]	[EB0200640000000000000000]
[EB]	Read Block	[0]	[EB0200640000000000000000]
[EB]	Write Block	[0]	[EB0200640000000000000000]
[EB]	Erase Block	[0]	[EB0200640000000000000000]
[EC]	Unstable Power Count	[0]	[EC0200646400000000000000]

```
(config) # license delete 1
```

2. Running

2.1 Introduction on my Ansible script

2.2.1 The Purpose

- It does MLNX-OS upgrade and downgrade continuously for replication test purpose.

While I am handling cases, sometimes, there are some customer complaints about “during OS/FW upgrade or downgrade, SSD gets stuck or errored block, Erase count increases”.

To verify that we TAC engineers sets up a lab. So to run it automatically via Ansible, we can save time and focus on the SSD data collected during OS upgrade/downgrade.

2.2.2 What it does

1). Doing OS upgrade -> downgrade -> upgrade.....

3.6.2102 -> 3.6.6162 -> 3.6.8010 -> 3.6.6162 -> 3.6.2102.....

2). During that, it writes current version, collects SSD status and write to a output.txt file

```
# print/write the current version
- name: print_current_version
  debug:
    msg: "Current version {{ current_version }}"
- name: Write Version
  local_action: shell echo "Current version {{ current_version }}" >> /root/output.txt

# Add license
- name: Add license
  onyx_config:
    lines:
      - license install XXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX
  register: add_license

# Get the SSD status
- name: _exec iSmart -d /dev/sda
```

```

    onyx_command:
      commands: _exec iSmart -d /dev/sda
      register: ssd_data

# Write the SSD status to output.txt file
- name: SSD status
  local_action: shell echo "{{ ssd_data.stdout }}" >> /root/output.txt

```

2.2.3 Exact steps

```

# Runing main task
- name: onyx_upgrade
~
~
tasks: ◀=== main task

# get current switch Onyx version with command "show version concise" and save it to a variable named curr_version
- name: get_current_switch_version

# save the onyx version as a fact named current_version
- name: set_fact_current_version

# print/write the current version
- name: print_current_version

# Add license
- name: Add license

# Get the SSD status
- name: _exec iSmart -d /dev/sda

# Write the SSD status to output.txt file
- name: SSD status

# run the delete images command
- name: delete images

# check if the img file exists in location
- name: check_source_file_existence

# Scp file copy
- name: copy_source_file_to_device

# Image install
- name: image install

# # image boot next
- name: image boot next

# # save config

```



```
- name: save config

# run the reload command
- name: reload

# use wait_for to wait for port 22 to be available on the switch after reboot
- name: wait for switch to reboot
```

2.2.4 How to run continuously.

Above “Exact steps” is 1 cycle. So it repeated only changing version info. So if you want to run many cycle, just copy & paste whole steps as many as you want to run.

2.2 How to run

```
[root@NVIDIA ~]# ansible-playbook OS_Up_downgrade.yml
```

➤ For debugging

```
[root@NVIDIA ~]# ansible-playbook OS_Up_downgrade.yml -vvvv
```

2.3 Running result : “Output.text”

```
[root@NVIDIA ~]# ansible-playbook OS_Up_downgrade.yml
```

```
[root@NVIDIA ~]# cat output.txt

Current version 3.6.2102

*****
* Innodisk iSMART V3.9.41                2018/05/25 *
*****
Model Name: Innodisk Corp. - mSATA 3ME
FW Version: S141119O
Serial Number: XXXXXXXXXXXXXXXXXXXX
Health: 0.00
Capacity: 14.914146 GB
P/E Cycle: 3000
Lifespan : 0 (Years : 0 Months : 0 Days : 0)
Write Protect: Disable
InnoRobust: Disable
```

ID	SMART Attributes	Value	Raw Value
[09]	Power On Hours	[53455]	[0900000000CFD00000000000]
[0C]	Power Cycle Count	[104]	[0C0000000068000000000000]
[AA]	Total Bad Block Count	[13]	[AA03006464000000D00000000]
[AD]	Erase Count Max.	[5861]	[AD12006464DE15E516000000]
[AD]	Erase Count Avg.	[5598]	[AD12006464DE15E516000000]
[C2]	Temperature	[0]	[000000000000000000000000]
[EB]	Later Bad Block	[0]	[EB0200640000000000000000]
[EB]	Read Block	[0]	[EB0200640000000000000000]
[EB]	Write Block	[0]	[EB0200640000000000000000]
[EB]	Erase Block	[0]	[EB0200640000000000000000]
[EC]	Unstable Power Count	[0]	[EC0200646400000000000000]

3. Possible script failures

3.1 Ansible & Python version mis-match.

➤ My lab environment. Please check and compare it with yours

- CentOS 8.0
- Ansible 2.9.24
- Python 3.6.8

3.2 License invalid.

➤ One of my script running is “Adding license & SSD status.”. Please modify/add license properly.

```
# Add license

- name: Add license
  onyx_config:
    lines:
      - license install XXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX
  register: add_license

# Get the SSD status

- name: _exec iSmart -d /dev/sda
  onyx_command:
    commands: _exec iSmart -d /dev/sda
  register: ssd_data
```

1). My script is written based on v3.6. So for this, license required to check SSD status.

2). If you want to skip to check SSD status, simply remove the above statement.

```
NVIDIA [standalone: master] # show licenses
License 1: XXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX_XXXXXX
Feature:   RESTRICTED_CMDS_GEN2
Description: Access to restricted system functionality
Valid:     yes
```

End date: 2023/07/13 (ok)
Tied to MAC addr: XX:XX:XX:XX:XX:XX (ok)
Active: yes ←=====

4. Reference

4.1 Ansible module for Mellanox

https://docs.ansible.com/ansible/latest/collections/mellanox/onyx/onyx_command_module.html

4.2 From NVIDIA Doc, how to

<https://docs.mellanox.com/display/ONYXv381208/Ansible>

<https://community.mellanox.com/s/article/Network-Automation-with-ZTP-and-Ansible-ONYX>

5. Errors

5.1 error messages running lower version Ansible.

```
TASK [copy_source_file_to_device]
*****
*****
task path: /root/hyungkwangc/OS_Up_Downgrade.yml:91
fatal: [l-csi-7700-0633.mtl.labs.mlnx]: FAILED! => {
  "changed": false,
  "destination": "/var/opt/tms/images/",
  "msg": "Exception received: command timeout triggered, timeout value is 30 secs.\nSee the timeout setting options in the
  Network Debug and Troubleshooting Guide."
}

PLAY RECAP
*****
*****
*****
l-csi-7700-0633.mtl.labs.mlnx : ok=10  changed=3  unreachable=0  failed=1  skipped=0  rescued=0  ignored=1

(env) [root@dell01 hyungkwangc]# ansible --version
ansible 2.9.21
  config file = /etc/ansible/ansible.cfg
  configured module search path = ['/root/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']
  ansible python module location = /opt/deepops/env/lib/python3.6/site-packages/ansible
  executable location = /opt/deepops/env/bin/ansible
  python version = 3.6.8 (default, Nov 16 2020, 16:55:22) [GCC 4.8.5 20150623 (Red Hat 4.8.5-44)]

(env) [root@dell01 hyungkwangc]# pip3 install ansible==2.9.24
Collecting ansible==2.9.24
  Downloading ansible-2.9.24.tar.gz (14.3 MB)
    14.3 MB 5.7 MB/s
Requirement already satisfied: jinja2 in /opt/deepops/env/lib/python3.6/site-packages (from ansible==2.9.24) (2.11.1)
Requirement already satisfied: PyYAML in /opt/deepops/env/lib64/python3.6/site-packages (from ansible==2.9.24) (5.4.1)
Requirement already satisfied: cryptography in /opt/deepops/env/lib64/python3.6/site-packages (from ansible==2.9.24) (3.4.8)
Requirement already satisfied: cffi>=1.12 in /opt/deepops/env/lib64/python3.6/site-packages (from cryptography->ansible==2.9.24) (1.14.6)
Requirement already satisfied: pycparser in /opt/deepops/env/lib/python3.6/site-packages (from cffi>=1.12->cryptography->ansible==2.9.24) (2.20)
Requirement already satisfied: MarkupSafe>=0.23 in /opt/deepops/env/lib64/python3.6/site-packages (from jinja2->ansible==2.9.24) (2.0.1)
Building wheels for collected packages: ansible
```

```

Building wheel for ansible (setup.py) ... done
Created wheel for ansible: filename=ansible-2.9.24-py3-none-any.whl size=16205052
sha256=8469484527c16441ce9842c705bea245db43361c25f85879ac10fbb31d47fb13
Stored in directory: /root/.cache/pip/wheels/fd/6f/e3/7fcd88f18084c9cf5fdaf9f9f0513ee486a90f97499caf5e6
Successfully built ansible
Installing collected packages: ansible
  Attempting uninstall: ansible
    Found existing installation: ansible 2.9.21
    Uninstalling ansible-2.9.21:
      Successfully uninstalled ansible-2.9.21
Successfully installed ansible-2.9.24 ←=====

(env) [root@dell01 hyungkwangc]# ansible --version
ansible 2.9.24
config file = /etc/ansible/ansible.cfg
configured module search path = ['/root/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']
ansible python module location = /opt/deepops/env/lib/python3.6/site-packages/ansible
executable location = /opt/deepops/env/bin/ansible
python version = 3.6.8 (default, Nov 16 2020, 16:55:22) [GCC 4.8.5 20150623 (Red Hat 4.8.5-44)]

```

5.2 Command timeout

```

TASK [copy_source_file_to_device]
*****
*****
task path: /root/hyungkwangc/OS_Up_Downgrade.yml:91
fatal: [l-csi-7700-0633.mtl.labs.mlnx]: FAILED! => {
  "changed": false,
  "destination": "/var/opt/tms/images/",
  "msg": "Exception received: command timeout triggered, timeout value is 30 secs.\nSee the timeout setting options in the
  Network Debug and Troubleshooting Guide."
}

PLAY RECAP
*****
*****
*****
l-csi-7700-0633.mtl.labs.mlnx : ok=10  changed=3  unreachable=0  failed=1  skipped=0  rescued=0  ignored=1

```

➤ How resolve

- Reference

https://docs.ansible.com/ansible/latest/network/user_guide/network_debug_troubleshooting.html#timeout-issues

- [Command timeout](#)

By default, `ANSIBLE_PERSISTENT_COMMAND_TIMEOUT` is set to 30 (seconds). Prior versions of Ansible had this value set to 10 seconds by default. You may see the following error if this value is too low:

```
2017-04-04 12:19:05,670 p=18591 u=fred | command timeout triggered, timeout value is 30 secs
```

Suggestions to resolve:

- Option 1 (Global command timeout setting): Increase value of command timeout in configuration file or by setting environment variable.

- `export ANSIBLE_PERSISTENT_COMMAND_TIMEOUT=60`

To make this a permanent change, add the following to your `ansible.cfg` file:

```
[persistent_connection]
command_timeout = 60
```

- Option 2 (Per task command timeout setting): Increase command timeout per task basis. All network modules support a timeout value that can be set on a per task basis. The timeout value controls the amount of time in seconds before the task will fail if the command has not returned.

For local connection type:

Suggestions to resolve:

```
- name: save running-config
  cisco.ios.ios_command:
    commands: copy running-config startup-config
    provider: "{{ cli }}"
    timeout: 30
```

Suggestions to resolve:

```
- name: save running-config
  cisco.ios.ios_command:
    commands: copy running-config startup-config
  vars:
    ansible_command_timeout: 60
```

Some operations take longer than the default 30 seconds to complete. One good example is saving the current running config on IOS devices to startup config. In this case, changing the timeout value from the default 30 seconds to 60 seconds will prevent the task from failing before the command completes successfully.

➤ What I did to resolve.

```
# The command timeout value defines the amount of time to wait for a command  
# or RPC call before timing out. The value for the command timeout must  
# be less than the value of the persistent connection idle timeout (connect_timeout)  
# The default value is 30 second.
```

```
command_timeout = 900 ←=== un-makred.
```